

AZHAR KHAN

Director / Head of · AI Operations, Implementation & Agentic Delivery

San Francisco, CA · 510-766-1727 · ap.apply@gmail.com · linkedin.com/in/azharpathan · azharkhann.github.io/portfolio

\$100M+ saved across a decade · 97.6% completion across 24.5M production runs · under 3¢ per business outcome

SUMMARY

Builds and runs the agentic orchestration layer a company runs on, across GTM, finance, and fraud. Keeps people on the decisions that carry risk, and hands everything else to governed AI agents. Three enterprises run on it today: a **top-5 US fintech** where an agent builds the fraud case and one analyst decides, a **national home-care platform** where nine jobs in 10,000 reach a person, and a **Microsoft-published utility operation** where one invoice in 100 does. Set the strategy for each, built each one, and still directs the team that runs them, reporting into each client's CTO or COO with full design-build-run authority.

CORE COMPETENCIES

Method: Process Discovery & Mining (Celonis) · Operating-Model & Process Reengineering · System / Data / Execution Gap Mapping · Digital-Twin Modeling · Golden-Record Master Data & Data Decoupling · Platform Selection & Build-vs-Buy · Service Operations (SLAs, Case Management, Escalation) · Client & Workforce Onboarding

Governance & Control Plane: AI Governance in Production · Agent Identity & Per-Action Entitlements · AI Gateway · Evals, Tracing & Observability · Drift & Hallucination Monitoring · Cost per Outcome (AI FinOps) · Enterprise MCP · LLMops / AgentOps

Execution Plane: Multi-Agent Orchestration · LLM & Agentic Frameworks · Context Engineering & Retrieval (RAG) · Intelligent Document Processing · Business Events · Connectivity & Integration (iPaaS) · Hybrid Deterministic + Agent + Human-in-the-Loop

Delivery & Change: Forward-Deployed Delivery · Embedded in Customer Operations · Program & Portfolio Delivery · Implementation & Rollout · Change Management & Adoption · Executive & Stakeholder Alignment · Benefits Realization

Stack: Vertex AI · Azure OpenAI · Azure AI Document Intelligence · LangGraph · BigQuery · Decagon · Sierra · Workato · UiPath · Appian · Datadog · Glean · Braintrust · Power Platform · Salesforce · Workday · NetSuite

EXPERIENCE

Concentrix

2019 – Present

Head of Agentic AI Engineering & Automation

San Francisco, CA

Built the automation Center of Excellence from 4 people to 30+ and now leads its 16-person Agentic AI Engineering & Automation practice, with \$10M+ in delivery revenue under management and 15+ engineers and architects hired across the US, Mexico, and India, running two to three enterprise transformations at once. The team is forward-deployed inside global enterprises, embedded in the client's operation to build and run these systems, and every benefit case clears a centralized CoE council before it is claimed. Builds missing infrastructure, makes an already-bought estate agent-ready, and has built the team that turns a broken operation into a running one, five times across five industries.

- ◆ **BotScope, the control plane:** built the fleet observability and governance layer in 2019 and productized it as a patented product, recognized by Forrester in 2021 (tracing, step-level failure attribution, MTTR, intelligent outage detection, proactive failure alerting, cross-bot telemetry). It now carries agent identity, per-action entitlements, a live execution graph, and token economics, so a single place answers who authorized an agent and what it cost.
- ◆ **CoE-in-a-Box:** productized the discovery-and-selection discipline. A five-day sprint replaced multi-month assessments, and phased first value ships at no new spend.

RECOGNITION · Microsoft-published case study (utility invoice processing) · Forrester recognition (BotScope, 2021) · UiPath AI10 Award · Patented product (BotScope)

PRODUCTION OPERATIONS

Governed agentic operations in production across financial services, healthcare, and utilities, where most enterprises never move past pilot.

Top-5 US Fintech · Agentic Fraud and Risk Operations · 2025 – Present

660K+ fraud cases a month · 7 lines of business · ~\$12M annual benefit · a human on every regulated decision

Seven financial-crimes-compliance (FCC) programs, scoped as seven separate builds. Investigators built each case by hand across six to eight systems, running a 21-step review of up to 150 data points that was wrong a third of the time. Set the strategy, and now directs the team running the platform in production across Asia and LATAM. After it proved out, leadership commissioned the global rollout; EMEA is now in design under the EU AI Act.

- ◆ **Platform selection, run as the buyer:** evaluated Decagon against Sierra head-to-head on live case volume and deployed each where it performed best; both now run in production across the seven lines of business.
- ◆ **Redesigned the work before automating it:** broke it into its atomic steps, re-sequenced disputes so denial parameters evaluate first, with both dispute service-level regimes (GNSR and GSND) mapped and their clocks encoded into the flow, and collapsed a 12-to-15-screen investigation into one workspace, cutting handle time by a third and lifting CSAT 13 points.
- ◆ **The agent builds the case; a human decides it:** an Automated Due Diligence agent pulls 129 data points from every system a case touches, scores them against the customer's baseline, checks SAR thresholds, and hands the analyst a decision-ready file with the report drafted. Every case is fully auditable.
- ◆ **Architecture and reliability:** an orchestrator runs specialized sub-agents (ingestion, investigation, reporting) over shared state, with gradient-boosted models scoring risk tiers 1 to 4 and hybrid retrieval assembling only the case's own evidence into context. Because the model is non-deterministic, the controls around it do the reliability work: evals gate every release on hallucination and drift, step-level tracing reconstructs any decision, guardrails contain prompt injection and enforce SAR thresholds, and PII stays tokenized.
- ◆ **Built once, reused seven times:** three checks (denylist, event-stream, UAR/SAR) run across all seven lines, and each new region deploys them directly, covering account takeover, disputes, elder financial abuse, identity verification, legal name change, and customer onboarding.

Global Bank · Fraud & Risk Operations Transformation · 2024 – 2025

Eight fraud journeys, application fraud to scams and card disputes · 188-FTE discovery baseline across 40 systems · designed tiered agentic model for 2.48M fraud calls a year

Ran a five-day discovery sprint across 188 FTE, 40 systems, and eight fraud journeys, spanning application and first-party fraud, scam claims, card disputes and chargebacks, fraud admin on SLA stop-clock mechanics (half the estate's volume), and securing funds. Then redesigned the operating model across four workstreams (journey-led transformation, contact strategy and correspondence, hyper-automation, and digital landscape assessment), taking 62 journey-system touchpoints to 3 intervention points. Directed it from discovery through operating-model design, and it now runs in production. The bank has since brought the practice back to design the future of its fraud and complaints operations.

- ◆ **Redesigned the case journey:** agentic case management collapses roughly 20 applications into a single shared case view, so a junior operator can work a case that used to need a specialist. Work auto-allocates and writes back downstream, letters issued automatically, with every investigation step ruled must, could, or won't per

decision class and product. GenAI summarizes the case, and a follow-up matrix automates twenty-plus post-decision actions per outcome, from CIFAS loads to fund-securing to bureau deletions. Stood up the Appian Center of Excellence against a £3.29M five-year savings plan.

- ◆ **Designed the tiered agentic model** for 2.48M fraud calls a year: automated resolution at level one, agent-assist at level two, and human investigators on romance scams, high-loss, and vulnerable-customer cases, with a phased path to case-management integration.

National In-Home Care Enterprise · Core Operating Platform, Governed Data & Agent Fleet · 2021 – Present

\$30M+ impact · 11 business functions on one governed platform · 97.6% completion across 24.5M runs over five years · nine jobs in 10,000 reach a person

Point-to-point manual handoffs across 35 disconnected systems, too fragmented to scale or carry AI. Set the strategy, mapped the operation, built one governed foundation in 2021, and has run it since. The whole company operates on it, across 12 US regions, under the CTO: 450 production workflows composing 3,650 callable services, 170M tasks a year. Agents consume it through an MCP layer, twenty use cases on top, no rebuild underneath.

- ◆ **Workforce, application to retained caregiver:** the pipeline runs 10,000+ hires a year across recruiting, background checks, HRIS, and payroll with zero orphaned accounts, spanning 23 lifecycle scenarios that took an 18% hiring surge without a rebuild; retention runs on the same platform, rewards flowing both directions, earned-wage access, and auto-generated caregiver bios, because in home care, supply is revenue capacity.
- ◆ **Revenue, first touch to collected cash:** the platform bills \$8.9M a week, runs the annual rate change as one orchestrated event, letters drafted for director approval and approved rates written back into the care platform, repricing the entire client book as a scheduled run; incoming payments match and apply to invoices automatically, conversions write back to the ad platforms, and month-end close is down 40%.
- ◆ **Finance, purchase order to payment:** runs multi-currency across subsidiaries, clears refunds form-to-approval-to-payment unattended, and generates contractor invoices directly from delivered care, suppliers synced from Workday into Coupa, so the platform writes the invoice instead of receiving it. One governed lake of 170B rows and a single golden record let finance, sales, and ops reconcile to one truth, closing \$1.33M in leakage. That record is decoupled from its source systems, so any one can be swapped without touching the process.
- ◆ **The back office runs on the same platform:** 121 automations carry warehousing, reverse-ETL, and master data on BigQuery; 25 carry support, cancellations, renewals, and NPS, with survey and call data landing beside tickets in the warehouse so churn surfaces early in the data; 22 carry ITSM, identity and access, and autonomous IT operations.
- ◆ **20 AI use cases, governed in production:** agent-callable skills through the enterprise MCP layer put 20 LLM and agentic use cases on the estate with no rewrites, including an agentic recruiting pipeline that replaced two hiring vendors. Decision rights are encoded per action, data stays access-scoped and CCPA-aligned, and every release passes eval gates with drift monitoring and a reconstructable audit trail.
- ◆ **The operation runs itself:** ~90% of a 30,000-ticket queue opens itself, nine jobs in 10,000 reach a person, and 146 core flows have run five years with zero production failures, absorbing nine acquisitions and three executive-sponsor changes as configuration. Four planes govern every job: per-action identity, 132 monitored tables, fail-closed recovery, design gates before any build. Each outcome runs under 3 cents; consumption fell 39% year over year as volume grew.
- ◆ **Readiness before agents:** scored 16 capabilities of a separately bought AI estate, ruled the nominated use cases unsafe to launch, and shipped a phased sequence instead, Phase 1 at no new spend. Roughly 40% of nominated use cases don't clear that review.

Global Utility-Bill Management Operation · Data Platform & Energy-Intelligence Transformation · Ongoing

700K invoices a month toward 1M · \$1B in payments · 180+ clients · 15,000+ vendors · Microsoft-published case study

A 700-person back-office operation ran the full utility-bill lifecycle by hand. The prior build had automated it with no documented process underneath, and it broke on real invoices. Owned the whole turnaround. Built a digital twin of the operation first, rebuilt the operation against it, and now operates it in production. The reconciled data is now the product; forecasting, rate optimization, benchmarking, and Scope 1 and 2 reporting all run off it. Microsoft published the result.

- ◆ **One reconciled record of spend, usage, and rate:** normalized thousands of billing sources across 15,000+ vendors and 300+ layouts into a single truth through a three-level charge dictionary, master definitions to vendor mappings to account charges. The twin runs as a live control loop, flagging deviations and naming the next automation off live data.
- ◆ **Capture, 65% to 99% accuracy:** three-path confidence routing (OCR for structured bills, trained models for known vendors, generative extraction for the unknowns) extracts to the meter and charge-line level, which is what lets the reconciled record carry forecasting and rate optimization, and cuts handle time per invoice from 8 minutes to under 1 and cost from \$0.58 to \$0.21, roughly \$3M a year, so one invoice in 100 now reaches a person.
- ◆ **Onboarding, 90 days to 30:** every month a client sits in onboarding leaves roughly \$800K of its contracted revenue idle. Re-engineered all ten phases and automated the two on the critical path. Account templates now pull 149 fields from invoice PDFs; the AP funding file builds in days.
- ◆ **Funding and reconciliation at scale:** absorbed a \$50M jump in monthly card volume with no added headcount, clearing ~1,900 payments a day. Historical loads gate at 99.5% reconciliation. The remittance audit now runs on three people.
- ◆ **Rebuild the operation and the team:** reshaped 290 document keyers into 65 pattern-builders and exception analysts, then designed nine specialized teams. The 225 freed roles absorb new clients and the climb from 700K toward 1M invoices a month. Made it agent-callable through the same MCP pattern proven at the care enterprise; it now extends to freight, manifests, and leases by configuration.

FOUNDATION: A DECADE OF PRODUCTION AUTOMATION AT ENTERPRISE SCALE

Delivered across Concentrix Catalyst (formerly PK) and earlier roles: a decade of automation programs totaling \$100M+ saved, built as automation, then bots, then agents, on the same foundations.

- ◆ **Fortune 200 insurer, claims:** brought AI-driven auditing to 540K claims a month, recovering \$7.4M at a 6.5x return where a 250-person manual operation had capped out.
- ◆ **Fortune 200 insurer, commissions:** rebuilt the function to run 1M invoices a year for 12,500 clients on a three-person team. UiPath AI10 Award.
- ◆ **Fortune 200 insurance services, finance operations:** automated the general-ledger process, cutting mean time to post the GL to 8 minutes and reducing days sales outstanding by a week.
- ◆ **Fortune 500 manufacturer:** built an enterprise automation program across four countries, returning \$6.5M+ and cutting the AP rebate cycle from 90 days to 2; a generative-AI pipeline (Stable Diffusion, ControlNet, LangGraph) turned engineering drawings into customer-ready imagery for 13,000+ products, cutting catalog production from weeks to hours, published as a Concentrix case study.
- ◆ **Fortune 50 telecom:** led a 250-automation program that returned \$18M; one redesign alone returned \$988K a year and freed 33 roles.

UiPath · Professional Services, Solution Engineering

2018 – 2019

- ◆ **BP and BPX Energy:** delivered 70+ automations worth \$20M, eliminated 100K+ hours, and owned the M&A and SAP S/4HANA automation roadmap.

ProKarma · Automation Consultant

2016 – 2018

- ◆ **Top US telecom:** built the enterprise automation governance, intake, and delivery frameworks; invoicing errors fell 85% and processing time fell 90%.

EDUCATION

M.S. Computer Science, San Francisco Bay University · B.Tech Computer Science, Vardhaman College of Engineering